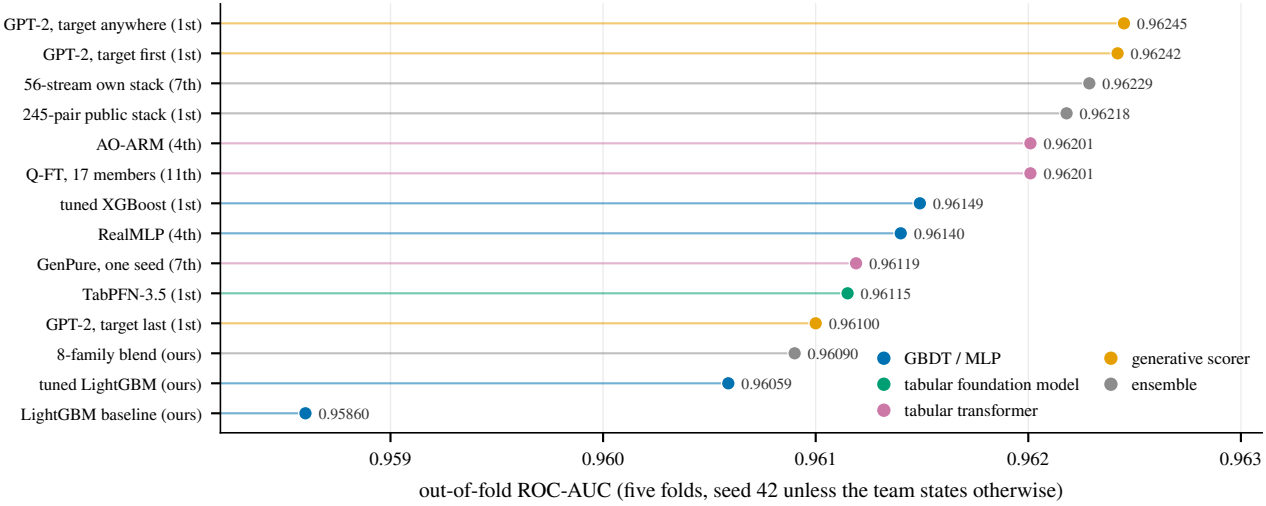
GPT-2, target anywhere (1st)
GPT-2, target first (1st)
56-stream own stack (7th)
245-pair public stack (1st)
AO-ARM (4th)
Q-FT, 17 members (11th)
tuned XGBoost (1st)
RealMLP (4th)
GenPure, one seed (7th)
TabPFN-3.5 (1st)
GPT-2, target last (1st)
8-family blend (ours)
tuned LightGBM (ours)
LightGBM baseline (ours)
0.96245
0.96242
0.96229
0.96218
0.96201
0.96201
0.96149
0.96140
0.96119
0.96115
0.96100
0.96090
0.96059
0.95860
GBDT / MLP
tabular foundation model
tabular transformer
generative scorer
ensemble
0.959
0.960
0.961
0.962
0.963
out-of-fold ROC-AUC (five folds, seed 42 unless the team states otherwise)